

An Introduction to **RcppEigen**

Douglas Bates

RcppEigen version 0.1.3 as of October 26, 2011

Abstract

The **RcppEigen** package provides access to the **Eigen** C++ template library from R. **Rcpp** classes and instantiations of the C++ templated functions **as** and **wrap** from **Rcpp** provide the “glue” for passing objects from R to C++ and back.

1 Introduction

As stated in the **Rcpp** vignette, “Extending **Rcpp**”

Rcpp facilitates data interchange between R and C++ through the templated functions **Rcpp::as** (for conversion of objects from R to C++) and **Rcpp::wrap** (for conversion from C++ to R).

The **RcppEigen** package provides the header files composing the **Eigen** C++ template library and implementations of **Rcpp::as** and **Rcpp::wrap** for the C++ classes defined in **Eigen**.

The **Eigen** classes themselves provide high-performance, versatile and comprehensive representations of dense and sparse matrices and vectors, as well as decompositions and other functions to be applied to these objects. In the next section we introduce some of these classes and show how to interface to them from R.

2 Eigen classes

Eigen (<http://eigen.tuxfamily.org>) is a C++ template library providing classes for many forms of matrices, vectors, arrays and decompositions. These classes are flexible and comprehensive allowing for both high performance and well structured code representing high-level operations. C++ code based on **Eigen** is often more like R code, working on the “whole object”, than compiled code in other languages where operations often must be coded in loops.

As in many C++ template libraries using template meta-programming (Abrahams and Gurtovoy, 2004), the templates themselves can be very complicated. However, **Eigen** provides typedef’s for common classes that correspond to R matrices and vectors, as shown in Table 1. We will use these typedef’s throughout this document.

The C++ classes shown in Table 1 are in the **Eigen** namespace, which means that they must be written as **Eigen::MatrixXd**. However, if we preface our use of these class names with a declaration like

```
using Eigen::MatrixXd;
```

we can use these names without the qualifier. I prefer this approach.

Table 1: Correspondence between R matrix and vector types and classes in the **Eigen** namespace.

R object type	Eigen class typedef
numeric matrix	MatrixXd
integer matrix	MatrixXi
complex matrix	MatrixXcd
numeric vector	VectorXd
integer vector	VectorXi
complex vector	VectorXcd
Matrix::dgCMatrix	SparseMatrix<double>

2.1 Mapped matrices in Eigen

Storage for the contents of matrices from the classes shown in Table 1 is allocated and controlled by the class constructors and destructors. Creating an instance of such a class from an R object involves copying its contents. An alternative is to have the contents of the R matrix or vector mapped to the contents of the object from the Eigen class. For dense matrices we use the Eigen templated class `Map`. For sparse matrices we use the Eigen templated class `MappedSparseMatrix`.

We must, of course, be careful not to modify the contents of the R object in the C++ code. A recommended practice is always to declare mapped objects as `const`.

2.2 Arrays in Eigen

For matrix and vector classes **Eigen** overloads the `'**'` operator to indicate matrix multiplication. Occasionally we want component-wise operations instead of matrix operations. The `Array` templated classes are used in **Eigen** for component-wise operations. Most often we use the `array` method for Matrix or Vector objects to create the array.

2.3 Structured matrices in Eigen

There are **Eigen** classes for matrices with special structure such as symmetric matrices, triangular matrices and banded matrices. For dense matrices, these special structures are described as “views”, meaning that the full dense matrix is stored but only part of the matrix is used in operations. For a symmetric matrix we need to specify whether the lower triangle or the upper triangle is to be used as the contents, with the other triangle defined by the implicit symmetry.

3 Some simple examples

C++ functions to perform simple operations on matrices or vectors can follow a pattern of:

1. Map the R objects passed as arguments into Eigen objects.
2. Create the result.
3. Return `Rcpp::wrap` applied to the result.

An idiom for the first step is

```
using Eigen::Map;
using Eigen::MatrixXd;
using Rcpp::as;

const Map<MatrixXd> A(as<Map<MatrixXd> >(AA));
```

where `AA` is the name of the R object (called an `SEXP` in C and C++) passed to the C++ function.

The `cxxfunction` from the **inline** package for R and its **RcppEigen** plugin provide a convenient way of developing and debugging the C++ code. For actual production code we generally incorporate the C++ source code files in a package and include the line `LinkingTo: Rcpp, RcppEigen` in the package's `DESCRIPTION` file.

The `cxxfunction` with the "Rcpp" or "RcppEigen" plugins has the `as` and `wrap` functions already defined as `Rcpp::as` and `Rcpp::wrap`. In the examples below we will omit these declarations. Do remember that you will need them in C++ source code for a package.

The first few examples are simply for illustration as the operations shown could be more effectively performed directly in R. We do compare the results from **Eigen** to those from the direct R results.

3.1 Transpose of an integer matrix

We create a simple matrix of integers

```
> (A <- matrix(1:6, ncol=2))

      [,1] [,2]
[1,]    1    4
[2,]    2    5
[3,]    3    6
```

```
> str(A)

int [1:3, 1:2] 1 2 3 4 5 6
```

and use the `transpose` method for the `Eigen::MatrixXi` class to return its transpose.

```
using Eigen::Map;
using Eigen::MatrixXi;
// Map the integer matrix AA from R
const Map<MatrixXi> A(as<Map<MatrixXi> >(AA));
// evaluate and return the transpose of A
const MatrixXi At(A.transpose());
return wrap(At);
```

```
> ftrans <- cxxfunction(signature(AA = "matrix"),
+                         paste(readLines("code.cpp"), collapse = "\n"),
+                         plugin = "RcppEigen")
> (At <- ftrans(A))

      [,1] [,2] [,3]
[1,]     1     2     3
[2,]     4     5     6

> stopifnot(all.equal(At, t(A)))
```

For numeric or integer matrices the `adjoint` method is equivalent to the `transpose` method. For complex matrices, the adjoint is the conjugate of the transpose. In keeping with the conventions in the `Eigen` documentation we prefer the name `adjoint` with numeric or integer matrices.

3.2 Products and cross-products

As mentioned in Sec. 2.2, the `'*'` operator performs matrix multiplication on `Matrix` or `Vector` objects.

```
using Eigen::Map;
using Eigen::MatrixXi;
const Map<MatrixXi> B(as<Map<MatrixXi> >(BB));
const Map<MatrixXi> C(as<Map<MatrixXi> >(CC));
return List::create(_["B %*% C"] = B * C,
                   _["crossprod(B, C)"] = B.adjoint() * C);
```

```
> fprod <- cxxfunction(signature(BB = "matrix", CC = "matrix"),
+                         paste(readLines("code.cpp"), collapse = "\n"), "RcppEigen")
> B <- matrix(1:4, ncol=2); C <- matrix(6:1, nrow=2)
> str(fp <- fprod(B, C))

List of 2
 $ B %*% C      : int [1:2, 1:3] 21 32 13 20 5 8
 $ crossprod(B, C): int [1:2, 1:3] 16 38 10 24 4 10

> stopifnot(all.equal(fp[[1]], B %*% C), all.equal(fp[[2]], crossprod(B, C)))
```

Notice that the `create` method for the `Rcpp` class `List` implicitly applies `Rcpp::wrap` to its arguments.

3.3 Crossproduct of a single matrix

As shown in the last example, the R function `crossprod` calculates the product of the transpose of its first argument with its second argument. The single argument form, `crossprod(X)`, evaluates $X'X$. We could, of course, calculate this product as

```
> t(X) %*% X
```

but `crossprod(X)` is roughly twice as fast because the result is known to be symmetric and only half the result needs to be calculated.. The function `tcrossprod` evaluates `crossprod(t(X))` without actually forming the transpose.

To express these calculations in Eigen we create a `SelfAdjointView`, which is a dense matrix of which only one triangle is used, the other triangle being inferred from the symmetry. (“self-adjoint” is equivalent to symmetric when applied to non-complex matrices.)

The `Eigen` class name is `SelfAdjointView`. The method for general matrices that produces such a view is called `selfadjointView`. Both require specification of either the `Lower` or `Upper` triangle.

For triangular matrices the class is `TriangularView` and the method is `triangularView`. The triangle can be specified as `Lower`, `UnitLower`, `StrictlyLower`, `Upper`, `UnitUpper` or `StrictlyUpper`.

For self-adjoint views the `rankUpdate` method adds a scalar multiple of $\mathbf{A}\mathbf{A}'$ to the current symmetric matrix. The scalar multiple defaults to 1.

```
using Eigen::Map;
using Eigen::MatrixXi;
using Eigen::Lower;

const Map<MatrixXi> A(as<Map<MatrixXi> >(AA));
const int m(A.rows()), n(A.cols());
MatrixXi AtA(MatrixXi(n, n).setZero().selfadjointView<Lower>().rankUpdate(A.adjoint()));
MatrixXi AAt(MatrixXi(m, m).setZero().selfadjointView<Lower>().rankUpdate(A));

return List::create(_["crossprod(A)"] = AtA,
                   _["tcrossprod(A)"] = AAt);
```

```
> fcprd <- cxxfunction(signature(AA = "matrix"),
+                       paste(readLines( "code.cpp" ), collapse = "\n"), "RcppEigen")
> str(crp <- fcprd(A))
```

List of 2

```
$ crossprod(A) : int [1:2, 1:2] 14 32 32 77
$ tcrossprod(A): int [1:3, 1:3] 17 22 27 22 29 36 27 36 45
```

```
> stopifnot(all.equal(crp[[1]], crossprod(A)), all.equal(crp[[2]], tcrossprod(A)))
```

To some, the expressions to construct `AtA` and `AAt` in that code fragment are compact and elegant. To others they are hopelessly confusing. If you find yourself in the latter group, you just need to read the expression left to right. So, for example, we construct `AAt` by creating a general integer matrix of size $m \times m$ (where \mathbf{A} is $m \times n$), ensure that all its elements are zero, regard it as a self-adjoint (i.e. symmetric) matrix using the elements in the lower triangle, then add $\mathbf{A}\mathbf{A}'$ to it and convert back to a general matrix form (i.e. the strict lower triangle is copied into the strict upper triangle).

For these products we could use either the lower triangle or the upper triangle as the result will be symmetrized before it is returned.

3.4 Cholesky decomposition of the crossprod

The Cholesky decomposition of the positive-definite, symmetric matrix, \mathbf{A} , can be written in several forms. Numerical analysts define the “LLt” form as the lower triangular matrix, \mathbf{L} , such that $\mathbf{A} = \mathbf{L}\mathbf{L}'$ and the “LDLt” form as a unit lower triangular matrix \mathbf{L} and a diagonal matrix \mathbf{D} with positive diagonal elements such that $\mathbf{A} = \mathbf{L}\mathbf{D}\mathbf{L}'$. Statisticians often write the decomposition as $\mathbf{A} = \mathbf{R}'\mathbf{R}$ where \mathbf{R} is an upper triangular matrix. Of course, this \mathbf{R} is simply the transpose of \mathbf{L} from the “LLt” form.

The templated `Eigen` classes for the LLt and LDLt forms are called `LLT` and `LDLT`. In general we would preserve the objects from these classes so that we could use them for solutions of linear systems. For illustration we simply return the matrix \mathbf{L} from the “LLt” form.

Because the Cholesky decomposition involves taking square roots we switch to numeric matrices.

```
> storage.mode(A) <- "double"
```

```

using Eigen::Map;
using Eigen::MatrixXd;
using Eigen::LLT;
using Eigen::Lower;

const Map<MatrixXd> A(as<Map<MatrixXd> >(AA));
const int n(A.cols());
const LLT<MatrixXd> llt(MatrixXd(n, n).setZero().
    selfadjointView<Lower>().rankUpdate(A.adjoint()));

return List::create(_["L"] = MatrixXd(llt.matrixL()),
    _["R"] = MatrixXd(llt.matrixU()));

```

```

> fchol <- cxxfunction(signature(AA = "matrix"),
+                       paste(readLines( "code.cpp" ), collapse = "\n"), "RcppEigen")
> (ll <- fchol(A))

$L
      [,1] [,2]
[1,] 3.741657 0.000000
[2,] 8.552360 1.963961

$R
      [,1] [,2]
[1,] 3.741657 8.552360
[2,] 0.000000 1.963961

> stopifnot(all.equal(ll[[2]], chol(crossprod(A))))

```

3.5 Determinant of the cross-product matrix

The “D-optimal” criterion for experimental design chooses the design that maximizes the determinant, $|\mathbf{X}'\mathbf{X}|$, for the $n \times p$ model matrix (or Jacobian matrix), \mathbf{X} . The determinant, $|\mathbf{L}|$, of the $p \times p$ lower Cholesky factor \mathbf{L} , defined so that $\mathbf{LL}' = \mathbf{X}'\mathbf{X}$, is the product of its diagonal elements. (This is true for any triangular matrix.) By the properties of determinants,

$$|\mathbf{X}'\mathbf{X}| = |\mathbf{LL}'| = |\mathbf{L}||\mathbf{L}'| = |\mathbf{L}|^2$$

Alternatively, if we use the “LDLt” decomposition, $\mathbf{LDL}' = \mathbf{X}'\mathbf{X}$ where \mathbf{L} is unit lower triangular and \mathbf{D} is diagonal then $|\mathbf{X}'\mathbf{X}|$ is the product of the diagonal elements of \mathbf{D} . Because we know that the diagonals of \mathbf{D} must be non-negative, we often evaluate the logarithm of the determinant as the sum of the logarithms of the diagonal elements of \mathbf{D} .

```

using Eigen::Lower;
using Eigen::Map;
using Eigen::MatrixXd;
using Eigen::VectorXd;

const Map<MatrixXd> A(as<Map<MatrixXd> >(AA));
const int n(A.cols());
const MatrixXd AtA(MatrixXd(n, n).setZero().
    selfadjointView<Lower>().rankUpdate(A.adjoint()));

const MatrixXd Lmat(AtA.llt().matrixL());
const double detL(Lmat.diagonal().prod());
const VectorXd Dvec(AtA.ldlt().vectorD());

return List::create(_["d1"] = detL * detL,
    _["d2"] = Dvec.prod(),
    _["ld"] = Dvec.array().log().sum());

```

```
> fdet <- cxxfunction(signature(AA = "matrix"),
+                       paste(readLines( "code.cpp" ), collapse = "\n"), "RcppEigen")
> unlist(ll <- fdet(A))

      d1      d2      ld
54.000000 54.000000 3.988984
```

Note the use of the `array()` method in the calculation of the log-determinant. Because the `log()` method applies to arrays, not to vectors or matrices, we must create an array from `Dvec` before applying the `log()` method.

4 Least squares solutions

A common operation in statistical computing is calculating a least squares solution, $\hat{\beta}$, defined as

$$\hat{\beta} = \arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2$$

where the model matrix, \mathbf{X} , is $n \times p$ ($n \geq p$) and \mathbf{y} is an n -dimensional response vector. There are several ways based on matrix decompositions, to determine such a solution. We have already seen two forms of the Cholesky decomposition: “LLt” and “LDLt”, that can be used to solve for $\hat{\beta}$. Other decompositions that can be used are the QR decomposition, with or without column pivoting, the singular value decomposition and the eigendecomposition of a symmetric matrix.

Determining a least squares solution is relatively straightforward. However, in statistical computing we often require a bit more information, such as the standard errors of the coefficient estimates. Calculating these involves evaluating the diagonal elements of $(\mathbf{X}'\mathbf{X})^{-1}$ and the residual sum of squares, $\|\mathbf{y} - \mathbf{X}\hat{\beta}\|^2$.

4.1 Least squares using the “LLt” Cholesky

```
using Eigen::LLT;
using Eigen::Lower;
using Eigen::Map;
using Eigen::MatrixXd;
using Eigen::VectorXd;

const Map<MatrixXd> X(as<Map<MatrixXd> >(XX));
const Map<VectorXd> y(as<Map<VectorXd> >(yy));
const int n(X.rows()), p(X.cols());
const LLT<MatrixXd> llt(MatrixXd(p, p).setZero().
    selfadjointView<Lower>().rankUpdate(X.adjoint()));
const VectorXd betahat(llt.solve(X.adjoint() * y));
const VectorXd fitted(X * betahat);
const VectorXd resid(y - fitted);
const int df(n - p);
const double s(resid.norm() / std::sqrt(double(df)));
const VectorXd se(s * llt.matrixL().solve(MatrixXd::Identity(p, p)).colwise().norm());
return List::create(_["coefficients"] = betahat,
    _["fitted.values"] = fitted,
    _["residuals"] = resid,
    _["s"] = s,
    _["df.residual"] = df,
    _["rank"] = p,
    _["Std. Error"] = se);
```

```
> llLS <- cxxfunction(signature(XX = "matrix", yy = "numeric"),
+                       paste(readLines( "code.cpp" ), collapse = "\n"), "RcppEigen")
> data(trees, package="datasets")
> str(lltFit <- with(trees, llLS(cbind(1, log(Girth)), log(Volume))))
```

```

List of 7
 $ coefficients : num [1:2] -2.35 2.2
 $ fitted.values: num [1:31] 2.3 2.38 2.43 2.82 2.86 ...
 $ residuals    : num [1:31] 0.0298 -0.0483 -0.1087 -0.0223 0.0727 ...
 $ s            : num 0.115
 $ df.residual  : int 29
 $ rank         : int 2
 $ Std. Error   : num [1:2] 0.2307 0.0898

> str(lmFit <- with(trees, lm.fit(cbind(1, log(Girth)), log(Volume))))

List of 8
 $ coefficients : Named num [1:2] -2.35 2.2
 ..- attr(*, "names")= chr [1:2] "x1" "x2"
 $ residuals    : num [1:31] 0.0298 -0.0483 -0.1087 -0.0223 0.0727 ...
 $ effects      : Named num [1:31] -18.2218 2.8152 -0.1029 -0.0223 0.0721 ...
 ..- attr(*, "names")= chr [1:31] "x1" "x2" "" "" ...
 $ rank         : int 2
 $ fitted.values: num [1:31] 2.3 2.38 2.43 2.82 2.86 ...
 $ assign       : NULL
 $ qr           :List of 5
 ..$ qr        : num [1:31, 1:2] -5.57 0.18 0.18 0.18 0.18 ...
 ..$ qraux: num [1:2] 1.18 1.26
 ..$ pivot: int [1:2] 1 2
 ..$ tol      : num 1e-07
 ..$ rank     : int 2
 ..- attr(*, "class")= chr "qr"
 $ df.residual : int 29

> for (nm in c("coefficients", "residuals", "fitted.values", "rank"))
+   stopifnot(all.equal(lltFit[[nm]], unname(lmFit[[nm]])))
> stopifnot(all.equal(lltFit[["Std. Error"]],
+   unname(coef(summary(lm(log(Volume) ~ log(Girth), trees)))[,2])))

```

There are several aspects of the C++ code worth mentioning. The `solve` method for the LLT object evaluates, in this case, $(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$ but without actually evaluating the inverse. The calculation of the residuals, $\mathbf{y} - \hat{\mathbf{y}}$, can be written, as in R, as `y - fitted`. (But note that **Eigen** classes do not have a “recycling rule as in R. That is, the two vector operands must have the same length.) The `norm` method evaluates the square root of the sum of squares of the elements of a vector. Although we don’t explicitly evaluate $(\mathbf{X}'\mathbf{X})^{-1}$ we do evaluate \mathbf{L}^{-1} to obtain the standard errors. Note also the use of the `colwise` method in the evaluation of the standard errors.

In the descriptions of other methods for solving least squares problems, much of the code parallels that shown here and we omit this. We show only the evaluation of the coefficients, the rank and the standard errors. Actually, we only calculate the standard errors up to the scalar multiple of s , the residual standard error, in these code fragments. The calculation of the residuals and s and the scaling of the coefficient standard errors is the same for all methods. (See the file `fastLm.cpp` of the **RcppEigen** source package for details.)

4.2 Least squares using the unpivoted QR decomposition

A QR decomposition has the form

$$\mathbf{X} = \mathbf{Q}\mathbf{R} = \mathbf{Q}_1\mathbf{R}_1$$

where \mathbf{Q} is an $n \times n$ orthogonal matrix, which means that $\mathbf{Q}'\mathbf{Q} = \mathbf{Q}\mathbf{Q}' = \mathbf{I}_n$, and \mathbf{R} is $n \times p$ and zero below the main diagonal. The $n \times p$ matrix \mathbf{Q}_1 is the first p columns of \mathbf{Q} and the $p \times p$ upper triangular matrix \mathbf{R}_1 is the top p rows of \mathbf{R} . There are three **Eigen** classes for the QR decomposition: **HouseholderQR** provides the basic QR decomposition using Householder transformations, **ColPivHouseholderQR** incorporates column pivots and **FullPivHouseholderQR** incorporates both row and column pivots.

For the unpivoted QR decomposition the code is of the form

```

using Eigen::HouseholderQR;

const HouseholderQR<MatrixXd> QR(X);
const VectorXd      betahat(QR.solve(y));
const VectorXd      fitted(X * betahat);
const int           df(n - p);
const VectorXd      se(QR.matrixQR().topRows(m_p).triangularView<Upper>().
                        solve(MatrixXd::Identity(m_p,m_p)).rowwise().norm());

```

We see that this code is very close to the code for the “LLt” case. In fact, if we had extracted the upper triangular factor (the `matrixU` method) from the LLT object, the code would be nearly identical.

4.3 Handling the rank-deficient case

One important consideration when determining least squares solutions is whether $\text{rank}(\mathbf{X})$ is p , a situation we describe by saying that \mathbf{X} has “full column rank”. When \mathbf{X} does not have full column rank we say it is “rank deficient”.

Although the theoretical rank of a matrix is well-defined, its evaluation in practice is not. At best we determine a computational rank according to some tolerance. We refer to decompositions that allow us to estimate the rank of the matrix in this way as “rank-revealing”.

Because the `model.matrix` function in R does a considerable amount of symbolic analysis behind the scenes, we usually end up with full-rank model matrices. The common cases of rank-deficiency, such as incorporating both a constant term and a full set of indicators columns for the levels of a factor, are eliminated. Other, more subtle, situations will not be detected at this stage, however. A simple example occurs when there is a “missing cell” in a two-way layout and the interaction of the two factors is included in the model.

```

> dd <- data.frame(f1 = gl(4, 6, labels = LETTERS[1:4]),
+                 f2 = gl(3, 2, labels = letters[1:3]))[-(7:8), ]
> xtabs(~ f2 + f1, dd)           # one missing cell

    f1
f2  A B C D
a  2 0 2 2
b  2 2 2 2
c  2 2 2 2

> mm <- model.matrix(~ f1 * f2, dd)
> kappa(mm)           # large condition number, indicating rank deficiency
[1] 4.309225e+16

> rcond(mm)           # alternative evaluation, the reciprocal condition number
[1] 2.320603e-17

> (c(rank=qr(mm)$rank, p=ncol(mm))) # rank as computed in R's qr function
rank      p
 11      12

> set.seed(1)
> dd$y <- mm %*% seq_len(ncol(mm)) + rnorm(nrow(mm), sd = 0.1)
>                                     # lm detects the rank deficiency
> fm1 <- lm(y ~ f1 * f2, dd)
> writeLines(capture.output(print(summary(fm1), signif.stars=FALSE))[9:22])

Coefficients: (1 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.97786    0.05816   16.81 3.41e-09
f1B          12.03807    0.08226  146.35 < 2e-16
f1C           3.11722    0.08226   37.90 5.22e-13
f1D           4.06852    0.08226   49.46 2.83e-14
f2b           5.06012    0.08226   61.52 2.59e-15

```


f2c	5.99759	0.08226	72.91	4.01e-16
f1B:f2b	-3.01476	0.11633	-25.92	3.27e-11
f1C:f2b	7.70300	0.11633	66.22	1.16e-15
f1D:f2b	8.96425	0.11633	77.06	< 2e-16
f1B:f2c	NA	NA	NA	NA
f1C:f2c	10.96133	0.11633	94.23	< 2e-16
f1D:f2c	12.04108	0.11633	103.51	< 2e-16

The `lm` function for fitting linear models in R uses a rank-revealing form of the QR decomposition. When the model matrix is determined to be rank deficient, according to the threshold used in R's QR decomposition, the model matrix is reduced to rank X columns by pivoting selected columns (those that are apparently linearly dependent on columns to their left) to the right hand side of the matrix. A solution for this reduced model matrix is determined and the coefficients and standard errors for the redundant columns are flagged as missing.

An alternative approach is to evaluate the “pseudo-inverse” of \mathbf{X} from the singular value decomposition (SVD) of \mathbf{X} or the eigendecomposition of $\mathbf{X}'\mathbf{X}$. The SVD is of the form

$$\mathbf{X} = \mathbf{U}\mathbf{D}\mathbf{V}' = \mathbf{U}_1\mathbf{D}_1\mathbf{V}'$$

where \mathbf{U} is an orthogonal $n \times n$ matrix and \mathbf{U}_1 is its leftmost p columns, \mathbf{D} is $n \times p$ and zero off the main diagonal so that \mathbf{D}_1 is a $p \times p$ diagonal matrix with non-decreasing positive diagonal elements, and \mathbf{V} is a $p \times p$ orthogonal matrix. The pseudo-inverse of \mathbf{D}_1 , written \mathbf{D}_1^+ is a $p \times p$ diagonal matrix whose first $r = \text{rank}(\mathbf{X})$ diagonal elements are the inverses of the corresponding diagonal elements of \mathbf{D}_1 and whose last $p - r$ diagonal elements are zero.

The tolerance for determining if an element of the diagonal of \mathbf{D} is considered to be (effectively) zero is a multiple of the largest singular value (i.e. the (1,1) element of \mathbf{D}).

We define a utility function, `Dplus` to return the pseudo-inverse as a diagonal matrix, given the singular values (the diagonal of \mathbf{D}) and the apparent rank.

```
using Eigen::DiagonalMatrix;
using Eigen::Dynamic;

inline DiagonalMatrix<double, Dynamic> Dplus(const ArrayXd& D, int r) {
    VectorXd Di(VectorXd::Constant(D.size(), 0.));
    Di.head(r) = D.head(r).inverse();
    return DiagonalMatrix<double, Dynamic>(Di);
}
```

4.4 Least squares with the SVD

With these definitions the code for least squares using the singular value decomposition can be written

```
using Eigen::JacobiSVD;

const JacobiSVD<MatrixXd> UDV(X.jacobiSvd(Eigen::ComputeThinU|Eigen::ComputeThinV));
const ArrayXd D(UDV.singularValues());
const int r((D > D[0] * threshold()).count());
const MatrixXd VDP(UDV.matrixV() * Dplus(D, r));
const VectorXd betahat(VDP * UDV.matrixU().adjoint() * y);
const int df(n - r);
const VectorXd se(s * VDP.rowwise().norm());
```

In the rank-deficient case this code will produce a complete set of coefficients and their standard errors. It is up to the user to note that the rank is less than p , the number of columns in \mathbf{X} and hence that the estimated coefficients are just one of an infinite number of coefficient vectors that produce the same fitted values. It happens that this solution is the minimum norm solution.

The interpretation of the standard errors from this code is also problematic when \mathbf{X} is rank-deficient.

4.5 Least squares with the eigendecomposition

The eigendecomposition of $\mathbf{X}'\mathbf{X}$ is defined as

$$\mathbf{X}'\mathbf{X} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}'$$

where \mathbf{V} , the matrix of eigenvectors, is a $p \times p$ orthogonal matrix and $\mathbf{\Lambda}$ is a $p \times p$ diagonal matrix with non-increasing, non-negative diagonal elements, called the eigenvalues of $\mathbf{X}'\mathbf{X}$. When the eigenvalues are distinct this \mathbf{V} is the same as that in the SVD. Also the eigenvalues of $\mathbf{X}'\mathbf{X}$ are the squares of the singular values of \mathbf{X} .

With these definitions we can adapt much of the code from the SVD method for the eigendecomposition.

```
using Eigen::SelfAdjointEigenSolver;

const SelfAdjointEigenSolver<MatrixXd>
    VLV(MatrixXd(p, p).setZero().selfadjointView<Lower>().rankUpdate(X.adjoint()));
const ArrayXd
    D(VLV.eigenvalues().array().sqrt());
const int
    r((D > D[0] * threshold()).count());
const MatrixXd
    VDP(UDV.matrixV() * Dplus(D, r));
const VectorXd
    betahat(VDP * VDP.adjoint() * X.adjoint() * y);
const int
    df(n - p);
const VectorXd
    se(s * VDP.rowwise().squaredNorm().array().sqrt());
```

4.6 Least squares with the column-pivoted QR decomposition

The column-pivoted QR decomposition provides results similar to those from R in both the full-rank and the rank-deficient cases. The decomposition is of the form

$$\mathbf{X}\mathbf{P} = \mathbf{Q}\mathbf{R} = \mathbf{Q}_1\mathbf{R}_1$$

where, as before, \mathbf{Q} is $n \times n$ and orthogonal and \mathbf{R} is $n \times p$ and upper triangular. The $p \times p$ matrix \mathbf{P} is a permutation matrix. That is, its columns are a permutation of the columns of \mathbf{I}_p . It serves to reorder the columns of \mathbf{X} so that the diagonal elements of \mathbf{R} are non-increasing in magnitude.

An instance of the class `Eigen::ColPivHouseholderQR` has a `rank` method returning the computational rank of the matrix. When \mathbf{X} is of full rank we can use essentially the same code as in the unpivoted decomposition except that we must reorder the standard errors. When \mathbf{X} is rank-deficient we evaluate the coefficients and standard errors for the leading r columns of $\mathbf{X}\mathbf{P}$ only.

In the rank-deficient case the straightforward calculation of the fitted values, as $\mathbf{X}\hat{\boldsymbol{\beta}}$ cannot be used. We could do some complicated rearrangement of the columns of \mathbf{X} and the coefficient estimates but it is conceptually (and computationally) easier to employ the relationship

$$\hat{\mathbf{y}} = \mathbf{Q}_1\mathbf{Q}_1'\mathbf{y} = \mathbf{Q} \begin{bmatrix} \mathbf{I}_r & \mathbf{00} & \mathbf{0} \end{bmatrix} \mathbf{Q}'\mathbf{y}$$

The vector $\mathbf{Q}'\mathbf{y}$ is called the “effects” vector in R.

```

using Eigen::ColPivHouseholderQR;
typedef ColPivHouseholderQR<MatrixXd>::PermutationType Permutation;

const ColPivHouseholderQR<MatrixXd> PQR(X);
const Permutation Pmat(PQR.colsPermutation());
const int r(PQR.rank());
VectorXd betahat, fitted, se;
if (r == X.cols()) { // full rank case
    betahat = PQR.solve(y);
    fitted = X * m_coef;
    se = Pmat * PQR.matrixQR().topRows(m_p).triangularView<Upper>().
        solve(MatrixXd::Identity(m_p, m_p)).rowwise().norm();
} else {
    MatrixXd Rinv(PQR.matrixQR().topLeftCorner(r, r).
        triangularView<Upper>().
        solve(MatrixXd::Identity(r, r)));
    VectorXd effects(PQR.householderQ().adjoint() * y);
    betahat.head(r) = Rinv * effects.head(r);
    betahat = Pmat * betahat;
    // create fitted values from effects
    // (cannot use X * m_coef when X is rank-deficient)
    effects.tail(X.rows() - r).setZero();
    fitted = PQR.householderQ() * effects;
    se.head(r) = Rinv.rowwise().norm();
    se = Pmat * se;
}

```

Just to check that this does indeed provide the desired answer

```
> print(summary(fmPQR <- fastLm(y ~ f1 * f2, dd)), signif.stars=FALSE)
```

Call:

```
fastLm.formula(formula = y ~ f1 * f2, data = dd)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.977859	0.058165	16.812	3.413e-09
f1B	12.038068	0.082258	146.346	< 2.2e-16
f1C	3.117222	0.082258	37.896	5.221e-13
f1D	4.068523	0.082258	49.461	2.833e-14
f2b	5.060123	0.082258	61.516	2.593e-15
f2c	5.997592	0.082258	72.912	4.015e-16
f1B:f2b	-3.014763	0.116330	-25.916	3.266e-11
f1C:f2b	7.702999	0.116330	66.217	1.156e-15
f1D:f2b	8.964251	0.116330	77.059	< 2.2e-16
f1B:f2c	NA	NA	NA	NA
f1C:f2c	10.961326	0.116330	94.226	< 2.2e-16
f1D:f2c	12.041081	0.116330	103.508	< 2.2e-16

Residual standard error: 0.2868 on 11 degrees of freedom

Multiple R-squared: 0.9999, Adjusted R-squared: 0.9999

```
> all.equal(coef(fm1), coef(fmPQR))
```

```
[1] TRUE
```

```
> all.equal(unname(fitted(fm1)), fitted(fmPQR))
```

```
[1] TRUE
```

```
> all.equal(unname(residuals(fm1)), residuals(fmPQR))
```

```
[1] TRUE
```

The rank-revealing SVD method produces the same fitted values but not the same coefficients.

```
> print(summary(fmSVD <- fastLm(y ~ f1 * f2, dd, method=4L)), signif.stars=FALSE)
```

Call:

```
fastLm.formula(formula = y ~ f1 * f2, data = dd, method = 4L)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.977859	0.058165	16.812	3.413e-09
f1B	7.020458	0.038777	181.049	< 2.2e-16
f1C	3.117222	0.082258	37.896	5.221e-13
f1D	4.068523	0.082258	49.461	2.833e-14
f2b	5.060123	0.082258	61.516	2.593e-15
f2c	5.997592	0.082258	72.912	4.015e-16
f1B:f2b	2.002847	0.061311	32.667	2.638e-12
f1C:f2b	7.702999	0.116330	66.217	1.156e-15
f1D:f2b	8.964251	0.116330	77.059	< 2.2e-16
f1B:f2c	5.017610	0.061311	81.838	< 2.2e-16
f1C:f2c	10.961326	0.116330	94.226	< 2.2e-16
f1D:f2c	12.041081	0.116330	103.508	< 2.2e-16

Residual standard error: 0.2868 on 11 degrees of freedom

Multiple R-squared: 0.9999, Adjusted R-squared: 0.9999

```
> all.equal(coef(fm1), coef(fmSVD))
```

```
[1] "'is.NA' value mismatch: 0 in current 1 in target"
```

```
> all.equal(unname(fitted(fm1)), fitted(fmSVD))
```

```
[1] TRUE
```

```
> all.equal(unname(residuals(fm1)), residuals(fmSVD))
```

```
[1] TRUE
```

For this example the method based on the eigendecomposition gets the rank wrong

4.7 Comparative speed

In the **RcppEigen** package the R function to fit linear models using the methods described above is called **fastLm**. The natural question to ask is, “Is it indeed fast to use these methods based on **Eigen**?”. We have provided a comparative benchmarking of these methods, plus the default method using R’s **lm** function and the **fastLm** implementations in the **RcppArmadillo** and **RcppGSL** packages, if they are installed. The benchmark is a file called **lmBenchmark.R** and it uses the **rbenchmark** package.

It can be run as

```
> source(system.file("examples", "lmBenchmark.R", package="RcppEigen"))
```

Results will vary according to the speed of the processor and the number of cores and the implementation of the BLAS (Basic Linear Algebra Subroutines) that is being used. (**Eigen** methods do not use the BLAS but the other methods do.)

Results obtained on a desktop computer, circa 2010, are shown in Table 2

These results indicate that methods based on forming and decomposing $\mathbf{X}'\mathbf{X}$, (i.e. LDLt, LLt and SymmEig) are considerably faster than the others. The SymmEig method, using a rank-revealing decomposition, would be preferred, although the LDLt method could probably be modified to be rank-revealing. Do bear in mind that the dimensions of the problem will influence the comparative results. Because there are 100,000 rows in \mathbf{X} methods that decompose the whole \mathbf{X} matrix (all the methods except those named above) will be at a disadvantage.

The pivoted QR method is 1.6 times faster than R’s **lm.fit** on this test and provides nearly the same information as **lm.fit**. Methods based on the singular value decomposition (SVD and GSL) are much slower but, as mentioned above, this is caused in part by \mathbf{X} having many more rows than columns. The GSL method from the GNU Scientific Library uses an older algorithm for the SVD and is clearly out of contention.

An SVD method using the Lapack SVD subroutine, **dgesv**, may be faster than the native **Eigen** implementation of the SVD, which is not a particularly fast method.

Table 2: **lmBenchmark** results on a desktop computer for the default size, $100,000 \times 40$, full-rank model matrix running 20 repetitions for each method. Times (Elapsed, User and Sys) are in seconds. The BLAS in use is a single-threaded version of Atlas (Automatically Tuned Linear Algebra System).

Method	Relative	Elapsed	User	Sys
LLt	1.000000	1.227	1.228	0.000
LDLt	1.037490	1.273	1.272	0.000
SymmEig	2.895681	3.553	2.972	0.572
QR	7.828036	9.605	8.968	0.620
PivQR	7.953545	9.759	9.120	0.624
arma	8.383048	10.286	10.277	0.000
lm.fit	13.782396	16.911	15.521	1.368
SVD	54.829666	67.276	66.321	0.912
GSL	157.531377	193.291	192.568	0.640

5 Delayed evaluation

A form of delayed evaluation is used in **Eigen**. That is, many operators and methods do not force the evaluation of the object but instead return an “expression object” that is evaluated when needed. As an example, even though we write the $\mathbf{X}'\mathbf{X}$ evaluation using `.rankUpdate(X.adjoint())` the `X.adjoint()` part is not evaluated immediately. The `.rankUpdate` method detects that it has been passed a matrix that is to be used in its transposed form and evaluates the update by taking inner products of columns of \mathbf{X} instead of rows of \mathbf{X}' .

Occasionally the method for `Rcpp::wrap` will not force an evaluation when it should. This is at least what Bill Venables calls an “infelicity” in the code, if not an outright bug. In the first example of the transpose of an integer matrix we assigned the transpose as a `MatrixXi` before returning it with `wrap`. The assignment forces the evaluation. If we skip this step we get an answer with the correct shape but the wrong contents.

```
using Eigen::Map;
using Eigen::MatrixXi;
const Map<MatrixXi> A(as<Map<MatrixXi> >(AA));
return wrap(A.transpose());
```

```
> Ai <- matrix(1:6, ncol=2L)
> ftrans2 <- cxxfunction(signature(AA = "matrix"),
+                          paste(readLines("code.cpp"), collapse = "\n"),
+                          plugin = "RcppEigen")
> (At <- ftrans2(Ai))

      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6

> all.equal(At, t(Ai))

[1] "Mean relative difference: 0.4285714"
```

Another recommended practice is to assign objects before wrapping them for return to R.

6 Sparse matrices

Eigen provides sparse matrix classes. An R object of class `dgCMatrix` (from the **Matrix** package) can be mapped as shown below.

```

using Eigen::Map;
using Eigen::MappedSparseMatrix;
using Eigen::SparseMatrix;
using Eigen::VectorXd;

const MappedSparseMatrix<double> A(as<MappedSparseMatrix<double> >(AA));
const Map<VectorXd> y(as<Map<VectorXd> >(yy));
const SparseMatrix<double> At(A.adjoint());
return List::create(_["At"] = At,
                    _["Aty"] = At * y);

```

```

> sparse1 <- cxxfunction(signature(AA = "dgCMatrix", yy = "numeric"),
+                         paste(readLines( "code.cpp" ), collapse = "\n"),
+                         plugin = "RcppEigen")
> data(KNex, package="Matrix")
> rr <- sparse1(KNex$mm, KNex$y)
> stopifnot(all.equal(rr$At, t(KNex$mm)),
+           all.equal(rr$Aty, as.vector(crossprod(KNex$mm, KNex$y))))

```

A sparse Cholesky decomposition is provided in **Eigen** as the **SimplicialCholesky** class. There are also linkages to the **CHOLMOD** code from the **Matrix** package. At present, both of these are regarded as experimental.

References

David Abrahams and Aleksey Gurtovoy. *C++ Template Metaprogramming: Concepts, Tools and Techniques from Boost and Beyond*. Addison-Wesley, Boston, 2004.